

The Blue Lotus Stress-Testing Engine: A Walk-Forward Out-of-Sample Validation

Sathvik Surapaneni

Blue Lotus Labs

sathviksurapaneni0@gmail.com

July 2026

Abstract

We report a walk-forward, out-of-sample validation of the Blue Lotus stress-testing engine (release v3.1) across a panel of 62 liquid instruments spanning seven asset classes and twelve annual test periods (2013–2024), yielding 744 independent evaluation cells. For each asset and test year the engine is estimated using only data available strictly before the test year, and its predicted one-year maximum-drawdown distribution is compared against the realized drawdown. This expanding-window design removes the single-period weakness of a one-shot hold-out: every year is a genuine out-of-sample test, and the panel spans multiple market regimes.

On the 496 cells drawn from non-crisis years, the engine’s 5% and 1% lower-tail bounds are breached at 4.0% and 1.0% respectively—rates for which a Kupiec proportion-of-failures test fails to reject nominal calibration ($p = 0.31$, $p = 0.99$; the test treats cells as independent, an assumption the panel only approximately satisfies). We do not claim uniform calibration of the whole distribution: a Kolmogorov–Smirnov test rejects PIT uniformity in every stratum, and in non-crisis years the departure is conservatism—the median probability-integral-transform (PIT) rank is 0.83, so realized drawdowns are typically milder than the engine’s central prediction, the prudent direction of error. Pooling all 744 cells the mean PIT rank is 0.53, but this near-central value masks two opposing effects—calm-year conservatism and crisis-year breaches—that average out. The residual tail miss is concentrated in the four crisis years (2015, 2018, 2020, 2022): there the 5% breach rate rises to 34.7%, degrading monotonically with severity and peaking in the March 2020 shock (67.7% breach, median PIT 0.09). Against four baselines (Gaussian, Student- t , and i.i.d. and block bootstraps) scored by the Continuous Ranked Probability Score, the engine is the only model with non-rejectable calm-market tail calibration and improves on every baseline’s forecast skill by 20–39% in non-crisis years, conceding sharpness only in crises—a deliberate trade the baselines make in reverse by being over-dispersed in all regimes. The engine therefore reproduces the frequency of ordinary tail losses faithfully out of sample and localizes—rather than conceals—the residual gap that a history-calibrated model cannot close.

Keywords: stress testing, walk-forward validation, out-of-sample backtest, maximum draw-down, Kupiec test, probability integral transform, calibration, tail risk.

This document validates a probability model. It does not constitute investment advice, and all figures are conditional on the modelling assumptions described in the companion engine paper.

Contents

1	Introduction and Motivation	4
2	The Engine Under Test (v3.1)	4
3	Walk-Forward Protocol	5
3.1	Expanding-window design	5
3.2	Coverage statistics	6
3.3	Universe	6
3.4	Benchmark models and scoring	6
4	Results	7
4.1	Headline calibration	7
4.2	Distributional uniformity	8
4.3	Calibration of predicted quantiles	8
4.4	Regime decomposition by year	9
4.5	Cross-asset behaviour	9
5	Comparison Against Baselines	10
6	Discussion	12
7	Limitations	14
8	Conclusion	15

1. Introduction and Motivation

A stress-testing engine is only as credible as the evidence that its predicted distributions actually contained what subsequently happened. A companion paper specifies the engine’s design and methodology in full and validates it on a single crisis-versus-calm event study. That design has two limitations this paper is written to remove.

First, an event study anchored to a fixed set of historical windows tests the engine at hand-chosen points in time. A more demanding and more honest protocol is *walk-forward* (expanding-window) validation: for each test period, the model is re-estimated using *only* the data available before that period, and evaluated on the period that follows. Repeating this across many consecutive periods yields a panel of strictly out-of-sample tests that (i) enforces train-before-test ordering everywhere, eliminating look-ahead, and (ii) spans every regime the sample contains, rather than a single hold-out window whose character—calm or turbulent—would otherwise dominate the verdict.

Second, the results reported here are produced by engine release **v3.1**, which corrects a set of systematic biases identified in an internal review of the prior release. Every correction moved reported risk in the conservative direction, so the calibration evidence below is measured on the corrected engine rather than its more optimistic predecessor. Section 2 summarises the changes.

The contribution is a single, large, reproducible calibration study: 744 out-of-sample cells across 62 instruments and 12 years, with proportion-of-failures and distributional-uniformity tests, stratified by market regime and by asset class.

2. The Engine Under Test (v3.1)

The numerical engine is unchanged in architecture from the companion paper: input processing and winsorization, a deterministic three-state volatility-regime model, Bayesian shrinkage of regime parameters, Generalized Pareto tail estimation via Peaks-over-Threshold Extreme Value Theory, regime-switching Monte Carlo path generation with a drawdown-conditional blend and a hard loss floor, an optional Esscher importance-resampling step, and bootstrapped confidence intervals. Release v3.1 corrects the following, all of which had biased reported risk optimistically:

1. **Tail estimation on raw losses.** The Generalized Pareto fit now uses the unwinsorized loss series; winsorizing before the fit had bounded the exceedances by construction and forced a spuriously short tail. Winsorization is retained only for bulk-moment estimation (regime detection and shrinkage).
2. **No path rejection.** An earlier step deleted the most extreme 1% of simulated paths before measurement; it is removed, so every path is retained.
3. **Consistent Esscher units.** The tilt’s mean and Expected-Shortfall constraints are now expressed in common daily-return units, so the change of measure can engage rather than

silently reverting to the identity.

4. **Row-aligned resampling.** Regime paths are reindexed together with return paths through any resampling step, so per-regime statistics pair returns with the correct regime sequence.
5. **Compounded drawdowns.** Drawdowns are computed on compounded wealth and are bounded below by -100% , on both the simulated and the realized side of every comparison.
6. **Stop-outs, not deletion.** A client-declared risk limit freezes a path at its loss rather than removing it from the sample.

For this study the engine is run with its documented production defaults: winsorization at the 1st/99th percentile for bulk moments; drawdown-conditioning thresholds at -5% (moderate) and -15% (severe); Extreme-Value level $\alpha = 0.05$; $N = 3,000$ Monte Carlo paths per cell; and a fixed random seed of 42. All drawdowns are one-year maximum drawdowns on compounded wealth.

3. Walk-Forward Protocol

3.1 Expanding-window design

Let an instrument have daily returns $\{r_t\}$ with dates $\{d_t\}$. For a test year Y , define the *training set* as every observation strictly before the first trading day of Y ,

$$\mathcal{T}_Y = \{r_t : d_t < 1 \text{ January } Y\},$$

and the *test set* as the returns realized during Y . The window is *expanding*: $\mathcal{T}_Y$ grows as Y advances, so each fit uses the maximum admissible history—appropriate for a tail model, which is data-hungry and has no reason to discard older crises. We require at least 504 training observations (≈ 2 years) and at least 200 realized test days per cell; cells failing either threshold are skipped.

The engine is estimated on $\mathcal{T}_Y$ and simulates a forward path ensemble of horizon equal to the number of trading days in Y , so the predicted maximum-drawdown distribution and the realized maximum drawdown span the same length. The realized maximum drawdown DD^{real} is computed on the raw (unwinsorized) test returns, in the same compounded-wealth convention as the simulation.

3.2 Coverage statistics

For each cell we record the realized maximum drawdown DD^{real} and the simulated distribution $\{DD^{(j)}\}_{j=1}^N$, and compute the *probability-integral-transform rank*

$$u = \frac{1}{N} \sum_{j=1}^N \mathbf{1}[DD^{(j)} \leq DD^{\text{real}}],$$

the simulated fraction at least as severe as reality. Under a correctly specified model $u \sim \text{Uniform}(0, 1)$. A $p5$ (*resp.* $p1$) *breach* occurs when DD^{real} is worse than the simulated 5th (*resp.* 1st) percentile. For x breaches in n cells at nominal rate p , the Kupiec proportion-of-failures statistic

$$\text{LR}_{\text{POF}} = -2 \ln \frac{(1-p)^{n-x} p^x}{(1-\hat{\pi})^{n-x} \hat{\pi}^x}, \quad \hat{\pi} = x/n,$$

is asymptotically χ_1^2 . We additionally test the full PIT sample for uniformity with a one-sample Kolmogorov–Smirnov statistic, and report the 90% band coverage—the fraction of cells with $DD^{\text{real}} \in [p5, p95]$.

3.3 Universe

The panel comprises 62 instruments chosen for liquidity and long history, grouped into seven asset classes: US equity beta (e.g. SPY, QQQ, IWM), sector funds (the nine SPDR sectors), international equity (EFA, EEM, and country funds), rates (TLT, IEF, SHY, AGG, TIP), credit (HYG, LQD, JNK), commodities (GLD, SLV, USO, DBC), and a set of long-history single-name equities. Crossed with the twelve test years, the grid yields 744 evaluation cells. Four years are designated *crisis* years for stratification—2015 (China/oil), 2018 (Volmageddon and the Q4 selloff), 2020 (COVID), and 2022 (the rate shock)—leaving eight non-crisis years.

3.4 Benchmark models and scoring

Calibration alone does not establish that the engine’s structure earns its complexity; a naive model could calibrate as well. We therefore evaluate four baselines on the *identical* cells, each fit on the same training window and each producing an $N = 3,000$ ensemble of one-year maximum drawdowns: (i) an i.i.d. Gaussian fit—the textbook Monte Carlo the engine is designed to improve on; (ii) an i.i.d. Student- t fit, a fat-tailed parametric competitor; (iii) a historical i.i.d. bootstrap of training returns; and (iv) a circular block bootstrap of block length 20, which preserves the volatility clustering the regime model targets and is the strongest baseline. Beyond the breach and PIT statistics, each cell is scored with the Continuous Ranked Probability Score,

$$\text{CRPS}(F, y) = \mathbb{E} |X - y| - \frac{1}{2} \mathbb{E} |X - X'|,$$

a strictly proper scoring rule for a probabilistic forecast of a scalar (with X, X' independent draws from the forecast ensemble F and y the realized drawdown; lower is better). We report mean CRPS, the engine’s CRPS skill score $1 - \overline{\text{CRPS}}_{\text{engine}} / \overline{\text{CRPS}}_{\text{base}}$ (positive favours the engine), the share of cells on which the engine’s CRPS is lower, and a paired Wilcoxon signed-rank test on the per-cell CRPS differences.

4. Results

4.1 Headline calibration

Table 1 reports the panel. The central result is a *tail-frequency* statement, and we state it as exactly that. On the 496 non-crisis cells the 5% bound is breached in 4.0% of cells and the 1% bound in 1.0%; a Kupiec proportion-of-failures test fails to reject nominal calibration at either level ($p = 0.31$ and $p = 0.99$). This is a statement about exception *rates*, not about the shape of the whole distribution: as Section 4.2 shows, the engine is mildly conservative through the body (non-crisis mean PIT rank 0.69), so it over-predicts risk in calm markets rather than under-predicting it—the prudent direction of error. Across the full panel the mean PIT rank is 0.53; this near-central average should not be read as distributional calibration, because it pools the calm-year conservatism above with the crisis-year breaches below (uniformity is rejected in every stratum, Section 4.2). The overall breach rates (14.3% at $p5$) are a blend dominated by the crisis subset, examined next.

Table 1: Walk-forward coverage: full panel, non-crisis years, and crisis years.

	All cells	Non-crisis	Crisis
Cells evaluated	744	496	248
$p5$ breach rate	14.3%	4.0%	34.7%
Kupiec p -value	$< 10^{-3}$	0.31	$< 10^{-3}$
$p1$ breach rate	5.7%	1.0%	14.9%
Kupiec p -value	$< 10^{-3}$	0.99	$< 10^{-3}$
90% band coverage	63.6%	63.5%	63.7%
Mean PIT rank	0.53	0.69	0.20

The 90% band coverage is near-identical across the two regimes (63.5% vs 63.7%) but for opposite reasons, and the distinction matters. In non-crisis years the band is missed almost exclusively on the *upper* side—realized drawdowns milder than the simulated 95th percentile—which is conservatism, not failure. In crisis years it is missed on the *lower* side, where realized losses fall below the simulated 5th percentile. The PIT distributions in Figure 1 make the asymmetry explicit: the non-crisis histogram leans gently right (mild over-prediction of risk), while the crisis histogram piles at the left edge (realized losses in the extreme predicted tail).




Figure 1: Realized-drawdown PIT ranks. Non-crisis years (left) track the uniform benchmark with a mild right lean; crisis years (right) concentrate at the left edge, where realized losses fall in the engine’s lower tail.

4.2 Distributional uniformity

The breach tests above concern exception *rates*; a stronger requirement is that the full PIT sample be uniform, and it is not. A one-sample Kolmogorov–Smirnov test rejects uniformity in every stratum: overall $D = 0.18$, non-crisis $D = 0.34$, crisis $D = 0.49$, all with $p < 10^{-3}$. These rejections have opposite sources and must not be collapsed into a single verdict. In non-crisis years the departure is *conservatism*—the median PIT rank is 0.83, so realized drawdowns are typically milder than the engine’s central prediction; the model over-states body risk, which for a risk system is the prudent failure mode. In crisis years the departure is the tail breach documented below (median PIT 0.09). The engine is therefore well calibrated in the specific sense the Kupiec tail test measures—the *frequency* of 5% and 1% exceedances in ordinary markets—but it is not a uniform generator of the entire drawdown distribution: it errs conservative through the body. We report this explicitly rather than let the favourable breach rates stand alone.

4.3 Calibration of predicted quantiles

Figure 2 plots empirical exceedance frequency against nominal lower-tail probability across the whole panel. Through the interior the empirical frequency runs at or below the nominal line—the graphical form of the body conservatism just described—and it lifts above the line only in the extreme lower tail, where the crisis-year breaches accumulate. The curve is thus the picture of a model whose exception *rates* are near nominal in the deep tail once crises are pooled in, while sitting conservative elsewhere.




Figure 2: Calibration curve. Empirical exceedance frequency versus nominal lower-tail probability, all 744 cells. The curve runs at or below the diagonal through the interior (body conservatism) and above it only in the extreme tail.

4.4 Regime decomposition by year

Table 2 resolves the panel by test year. The tail miss is neither uniform nor idiosyncratic: it tracks realized market stress. Six of the eight non-crisis years breach at or below the nominal 5% rate, several well below; the two exceptions—2013 (8.1%) and 2023 (11.3%, the year of the regional-bank episode)—exceed it modestly but remain far short of any crisis year. Among the crisis years, 2015 is fully contained (3.2% breach, 91.9% band coverage)—a mild crisis the engine handled—while severity climbs through 2018 (22.6%) and 2022 (45.2%) to 2020 (67.7%, mean PIT 0.06), the fastest broad equity drawdown in the sample. The monotone relationship between realized severity and breach rate is the central qualitative finding: the engine does not miss at random; it misses in proportion to how much history failed to anticipate the future.

Figure 3 shows the same decomposition graphically: eight bars at or below the nominal 5% line, and four crisis bars rising above it in proportion to severity.

4.5 Cross-asset behaviour

Table 3 pools cells by asset class. The tail miss concentrates in the exposures that carry the most violent crisis drawdowns—US equity beta (16.7% breach) and, notably, rates (18.3%), the latter driven by the historic 2022 Treasury drawdown that no pre-2022 window anticipated. Credit (8.3%) and the sector funds (10.2%) sit closer to nominal. Mean PIT ranks cluster around 0.5

Table 2: Coverage by test year ($n = 62$ each). Crisis years in the lower block.

Year	$p5$ breach	$p1$ breach	90% cov.	Mean PIT
2013	8.1%	1.6%	37.1%	0.76
2014	3.2%	0.0%	61.3%	0.74
2016	1.6%	0.0%	88.7%	0.62
2017	1.6%	0.0%	35.5%	0.86
2019	3.2%	0.0%	54.8%	0.77
2021	0.0%	0.0%	72.6%	0.65
2023	11.3%	3.2%	87.1%	0.44
2024	3.2%	3.2%	71.0%	0.66
2015	3.2%	1.6%	91.9%	0.42
2018	22.6%	3.2%	75.8%	0.24
2022	45.2%	16.1%	54.8%	0.08
2020	67.7%	38.7%	32.3%	0.06

across classes (from 0.43 for rates to 0.64 for US equity), so the central prediction is reasonably located across the cross-section and the dispersion in breach rates is a tail phenomenon rather than a directional bias in any one sleeve.

Table 3: Coverage by asset class, pooled over all test years.

Asset class	Cells	$p5$ breach	90% cov.	Mean PIT
Credit	36	8.3%	66.7%	0.53
Sector	108	10.2%	60.2%	0.58
Intl. equity	84	13.1%	65.5%	0.55
Commodity	48	14.6%	68.8%	0.48
Single name	312	15.1%	67.0%	0.49
US equity	96	16.7%	49.0%	0.64
Rates	60	18.3%	66.7%	0.43

Figure 4 plots the predicted 5th-percentile drawdown against the realized drawdown for every cell. The mass lies on or inside the 45° line; the breaches—points below it—are overwhelmingly the crisis-year cells in the lower-left, and the most extreme are dominated by 2020 and 2022 exposures such as energy (-61% realized) and long Treasuries (-35% realized).

5. Comparison Against Baselines

Table 4 places the engine against the four baselines of Section 3.4. The result is regime-dependent, and stating it precisely is the point of the exercise.

In non-crisis markets—where a risk system operates the great majority of the time—the engine dominates on both axes at once. It is the *only* model whose tail calibration is not




Figure 3: Out-of-sample p_5 breach rate by test year. Crisis years shaded; the horizontal line marks the nominal 5% rate.

rejected (4.0% p_5 breach, Kupiec $p = 0.31$); every baseline breaches at 0.6–0.8%, six to eight times below nominal, which means they systematically over-state tail risk and would raise near-constant false alarms. That same over-dispersion costs them sharpness: the engine’s mean CRPS is 20–39% lower than every baseline (paired Wilcoxon $p < 10^{-3}$ throughout), and it wins on 71–82% of calm cells. Pooled over the full panel the engine still posts the lowest CRPS, beating the Gaussian, Student- t , and i.i.d. bootstrap by 7–14% and edging the block bootstrap by 0.6% (a statistical tie, Wilcoxon $p = 0.04$).

In crisis years the ranking inverts: the engine’s CRPS is 27–38% *worse* than the baselines, and it wins only about a quarter of crisis cells. This is both honest and informative. The baselines score better in crises not because they are calibrated—none is; all four fail the Kupiec test at $p < 10^{-3}$, with p_5 breach rates of 21–22%—but because they are uniformly over-dispersed *everywhere*, so a permanently-too-wide distribution happens to sit closer to a realized extreme. The very sharpness that wins the engine the calm regime is what costs it in the tail. A model that is vague in all weather looks better in a storm and worse the other four-fifths of the year; the engine makes the opposite trade deliberately, and the operator-configurable crisis overlay described in the companion engine paper is the lever for widening the distribution when sizing for systemic scenarios.

Figure 5 shows the CRPS decomposition and Figure 6 the calm-market calibration gap. The conclusion a technical reader should draw is specific. The regime, Extreme-Value, and shrinkage machinery is not decoration: it is the only configuration in this comparison that delivers *both* calibrated and sharp calm-market drawdown forecasts, which naive parametric or historical simulation does not, and it does so at a cost in crisis sharpness that is explicit and addressable rather than concealed. A buyer evaluating the engine on the regime it will spend most of its life in is evaluating its strongest case; a buyer evaluating it on crises alone would do better with a




Figure 4: Predicted 5th-percentile drawdown versus realized annual maximum drawdown, one point per cell. Points below $y = x$ are tail breaches; the lower-left cluster is the 2020 and 2022 crisis cells.

deliberately vague model, and should instead engage the crisis overlay.

6. Discussion

The walk-forward panel delivers a sharper version of the message from the earlier event study, and it does so on the corrected engine. On ordinary market conditions—the regime in which a risk system spends the overwhelming majority of its life—a Kupiec test fails to reject nominal 5% and 1% tail-breach rates out of sample across 496 asset-years, with the body of the distribution biased toward caution. We are careful about the strength of that claim: it concerns exception frequencies, not the full distribution (which is not uniform, Section 4.2), and the test assumes independent cells, which the panel only approximately satisfies (Section 7). Within those bounds it is the precondition for taking any crisis result seriously: a model that over-fired in calm markets would provide no standing to interpret its behaviour in storms.

The residual miss is real, and the study measures rather than hides it. It is concentrated in the crisis years, monotone in realized severity, and concentrated in the exposures that carry systemic tail risk. The natural explanation is not that the model had seen no crisis: price history begins in 2004, so every training window for the 2020 test cells spans twelve to sixteen years and *includes the 2008 global financial crisis* without exception. The engine under-covered 2020 despite having been fit on 2008. The operative fact is that a tail calibrated to one systemic

Table 4: Benchmark comparison. Mean CRPS (lower is better) by regime; non-crisis tail calibration; and the engine’s CRPS skill score against each baseline on non-crisis cells (positive favours the engine). $\text{CRPS} \times 10^2$.

Model	CRPS calm	CRPS crisis	CRPS all	Calm $p5$ breach	Engine skill (calm)
Blue Lotus	4.20	9.32	5.91	4.0% ($p = 0.31$)	—
Gaussian	6.89	6.77	6.85	0.8% ($p < 10^{-3}$)	+39%
Student- t	5.94	7.20	6.36	0.8% ($p < 10^{-3}$)	+29%
Hist. bootstrap	6.35	6.88	6.53	0.8% ($p < 10^{-3}$)	+34%
Block bootstrap	5.24	7.34	5.94	0.6% ($p < 10^{-3}$)	+20%




Figure 5: Mean CRPS by model and regime (lower is better). Blue Lotus is best in non-crisis years and pooled, and worst in crises—the deliberate sharpness-versus-robustness trade discussed in the text.

episode need not envelope the next: the March 2020 drawdown’s velocity and cross-asset breadth differed from the 2008 shape the model had learned, in exactly the dimensions that drive a one-year maximum drawdown. A history-calibrated tail is bounded by the worst *shape* it has observed, and 2020 exceeded 2008 on that axis. The design response is already in the engine—the crisis-regime structure, the drawdown-conditional blend, and the operator-configurable loss floor exist to extend the simulated distribution into the region this panel shows lies beyond historical reach. The honest deployment is to use the engine as calibrated here for ordinary risk measurement, and to engage the crisis overlay when sizing for systemic scenarios.

It is worth recording what the v3.1 corrections bought, without overstating it. The prior release’s optimistic biases—winsorized tails, deletion of extreme paths, additive drawdowns—would each have pushed realized outcomes further into the predicted tail, lowering PIT ranks and inflating breach rates; measuring the corrected engine is what allows the calm-year breach frequencies to land at nominal rather than beyond it. It does not make the distribution uniform:




Figure 6: Non-crisis p_5 breach rate by model. Only Blue Lotus sits near the nominal 5% line; every baseline is far below it, over-stating tail risk in calm markets.

the corrected engine is conservative through the body, which the corrections did not, and were not intended to, remove.

7. Limitations

1. **Cross-sectional dependence.** The 62 instruments in a given year share a common macro environment, so the effective number of independent observations within a year is smaller than the cell count. The Kupiec and Kolmogorov–Smirnov statistics reported here assume independent cells; under the dependence present in the panel their p -values overstate precision, and the “fail to reject” verdict on non-crisis calibration should be read as suggestive rather than a formal size-controlled test. A year-clustered or block-bootstrap treatment is the correct next step and is left to future work. This dependence is also why the panel’s power derives chiefly from its *temporal* span across regimes, not from cross-sectional breadth.
2. **History depth.** Price data begin in 2004. Every training window therefore contains the 2008 crisis, but none contains a pre-2004 regime; the earliest test years (from 2013) are fit on roughly nine years of data, and the tail estimates rest on at most ~ 16 years of exceedances. No 1987, 2000, or earlier stress informs the fits, and the deep-tail parameters are correspondingly data-limited.
3. **Annual horizon.** Each cell evaluates a one-year maximum drawdown. Shorter or path-dependent horizons would probe different aspects of the predicted distribution and are left to future work.
4. **Survivorship and vendor.** The universe is composed of instruments that survived to the present, and all prices come from a single adjusted-close source; both would shift individual cells modestly without altering the aggregate picture.

5. **Crisis-year labelling.** The calm/crisis split is a coarse stratification for interpretation; the continuous relationship in Table 2 does not depend on it.

8. Conclusion

Under a strict expanding-window protocol spanning 62 instruments, seven asset classes, and twelve years—744 out-of-sample cells in total—engine release v3.1 reproduces the *frequency* of ordinary tail losses on non-crisis markets: a Kupiec test fails to reject nominal 5% and 1% breach rates (subject to the cross-sectional-dependence caveat of Section 7), while the engine errs conservative through the body of the distribution, which a uniformity test confirms is not itself uniform. Its behaviour under the four crisis years of the period is fully characterised: coverage degrades monotonically with realized severity and breaks down only at the largest systemic shocks, exactly where a tail calibrated to prior crises must. Against four naive baselines scored by a strictly proper rule, it is the sole model delivering *calibrated and sharp* calm-market forecasts—improving CRPS by 20–39% in the regime where a risk desk operates—while conceding crisis sharpness by construction, a trade the crisis overlay exists to reverse on demand. The standard adopted throughout is the one that matters: quantify the residual gap out of sample, state precisely what was and was not demonstrated, and report both.

Reproducibility. The panel regenerates from `backtest/walk_forward.py` (evaluation) and `backtest/wf_analyze.py` (aggregation and figures) with fixed seed 42; per-cell output is written to `backtest/wf_results/cells.json`.